

Predicting Startup Success Using Investment and Growth Data

March 28, 2025

1 Objectives of this program

- To explore the factors that can improve the **Growth Rate (%)** of a Startup business
- To explore the relationship between the factors that can increase the **Growth Rate** of a startup
- To create a machine learning model that can predict the **Growth Rate** of a startup based off of those factors
- To provide solutions to make a successful startup

```
[1]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import warnings

warnings.filterwarnings("ignore")
```

2 Exploratory Data Analysis

Loading the dataset

```
[2]: file_path = "startup_growth_investment_data.csv"
df = pd.read_csv(file_path)

# Remove extra white space from columns
df.columns = df.columns.str.strip()

# Drop the useless column
df = df.drop("Startup Name", axis=1)
```

```
[3]: # Display the first 5 rows of data
df.head()
```

	Industry	Funding Rounds	Investment Amount (USD)	Valuation (USD)	\
0	Blockchain	8	1.335166e+09	6.621448e+09	
1	SaaS	2	2.781498e+09	8.363214e+09	
2	EdTech	10	3.309032e+09	1.548270e+10	

```

3      Fintech          5           4.050196e+09    1.268253e+10
4      EdTech          9           1.645080e+09    6.887966e+09

```

	Number of Investors	Country	Year Founded	Growth Rate (%)
0	50	Germany	2012	77.10
1	36	UK	2006	105.52
2	39	Singapore	2016	190.47
3	44	France	2021	9.44
4	48	India	2011	192.00

[4]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 8 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Industry          5000 non-null   object 
 1   Funding Rounds    5000 non-null   int64  
 2   Investment Amount (USD) 5000 non-null   float64
 3   Valuation (USD)     5000 non-null   float64
 4   Number of Investors 5000 non-null   int64  
 5   Country            5000 non-null   object 
 6   Year Founded       5000 non-null   int64  
 7   Growth Rate (%)    5000 non-null   float64
dtypes: float64(3), int64(3), object(2)
memory usage: 312.6+ KB

```

[5]: # Statistical Summary
df.describe()

	Funding Rounds	Investment Amount (USD)	Valuation (USD)	\
count	5000.000000	5.000000e+03	5.000000e+03	
mean	5.491600	2.455567e+09	7.971059e+09	
std	2.913353	1.423787e+09	5.479487e+09	
min	1.000000	1.102610e+06	1.722547e+06	
25%	3.000000	1.221506e+09	3.598305e+09	
50%	6.000000	2.460634e+09	7.002304e+09	
75%	8.000000	3.639951e+09	1.147676e+10	
max	10.000000	4.999544e+09	2.470906e+10	

	Number of Investors	Year Founded	Growth Rate (%)
count	5000.000000	5000.000000	5000.000000
mean	25.542000	2011.544000	102.091732
std	14.271838	6.885285	56.179781
min	1.000000	2000.000000	5.060000
25%	13.000000	2006.000000	52.815000
50%	25.000000	2012.000000	102.215000

```
75%           38.000000  2018.000000      150.580000
max          50.000000  2023.000000     199.970000
```

```
[6]: # Checking column names
df.columns
```

```
[6]: Index(['Industry', 'Funding Rounds', 'Investment Amount (USD)', 'Valuation (USD)', 'Number of Investors', 'Country', 'Year Founded', 'Growth Rate (%)'],
       dtype='object')
```

```
[7]: # Checking data types
df.dtypes
```

```
[7]: Industry          object
Funding Rounds      int64
Investment Amount (USD) float64
Valuation (USD)      float64
Number of Investors  int64
Country             object
Year Founded        int64
Growth Rate (%)     float64
dtype: object
```

```
[8]: # Checking for missing values
missing_values = df.isnull().sum()
missing_values
```

```
[8]: Industry          0
Funding Rounds      0
Investment Amount (USD) 0
Valuation (USD)      0
Number of Investors  0
Country             0
Year Founded        0
Growth Rate (%)     0
dtype: int64
```

```
[9]: # Checking for duplicated values
duplicated_values = df.duplicated().sum()
duplicated_values
```

```
[9]: 0
```

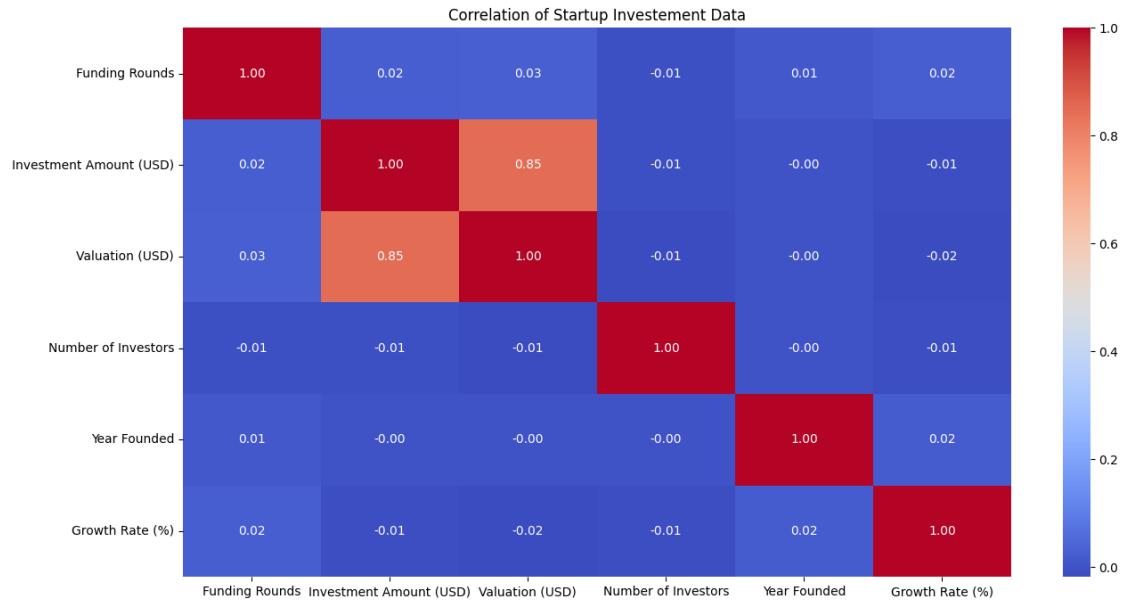
```
[10]: numerical_df = df.select_dtypes(include=["int64", "float64"])
```

```
# Compute and show the heatmap
plt.figure(figsize=(15, 8))
```

```

sns.heatmap(numerical_df.corr(), annot=True, fmt='.2f', cmap='coolwarm')
plt.title("Correlation of Startup Investement Data")
plt.show()

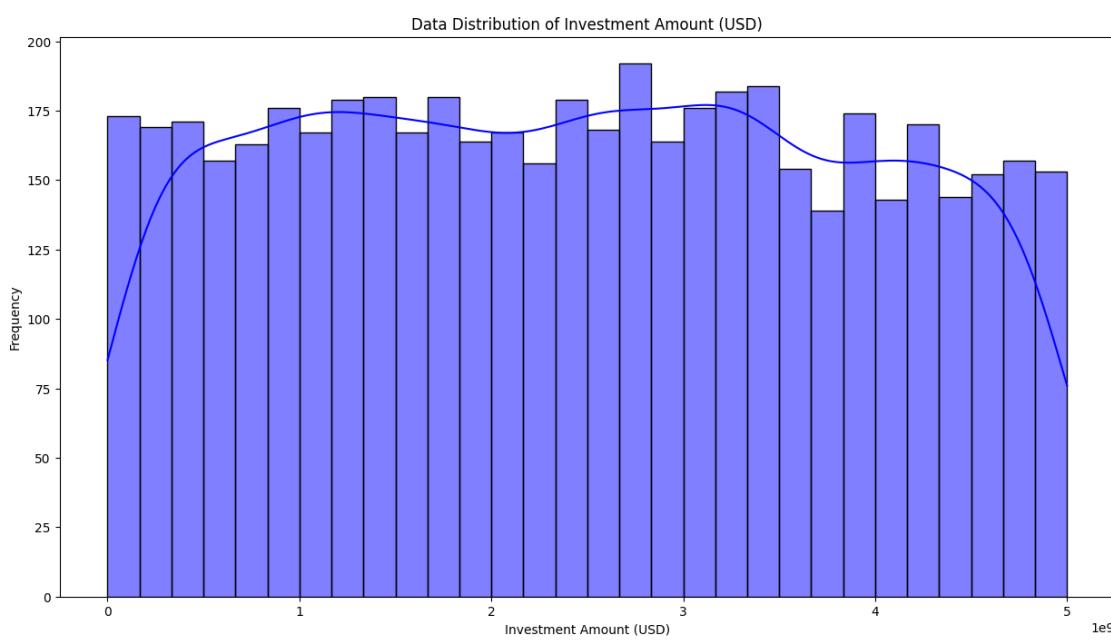
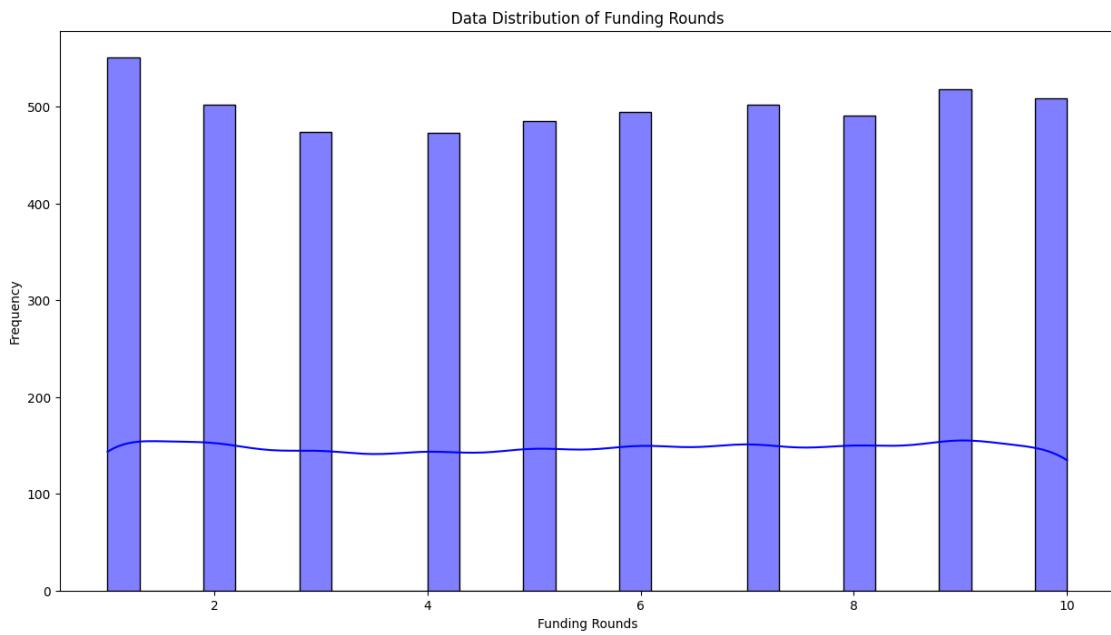
```

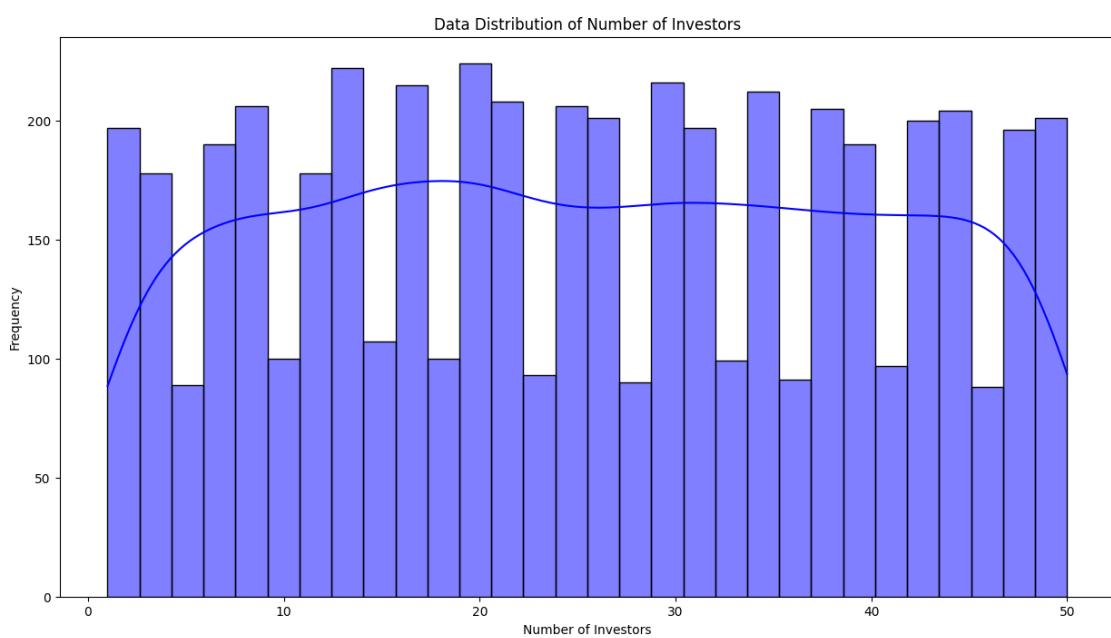
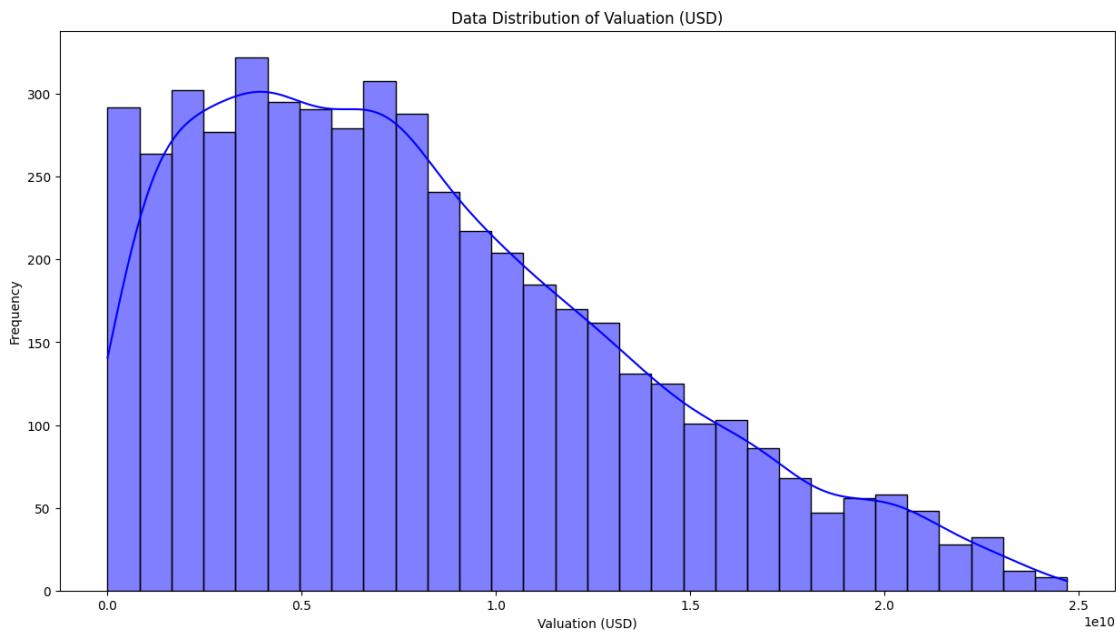


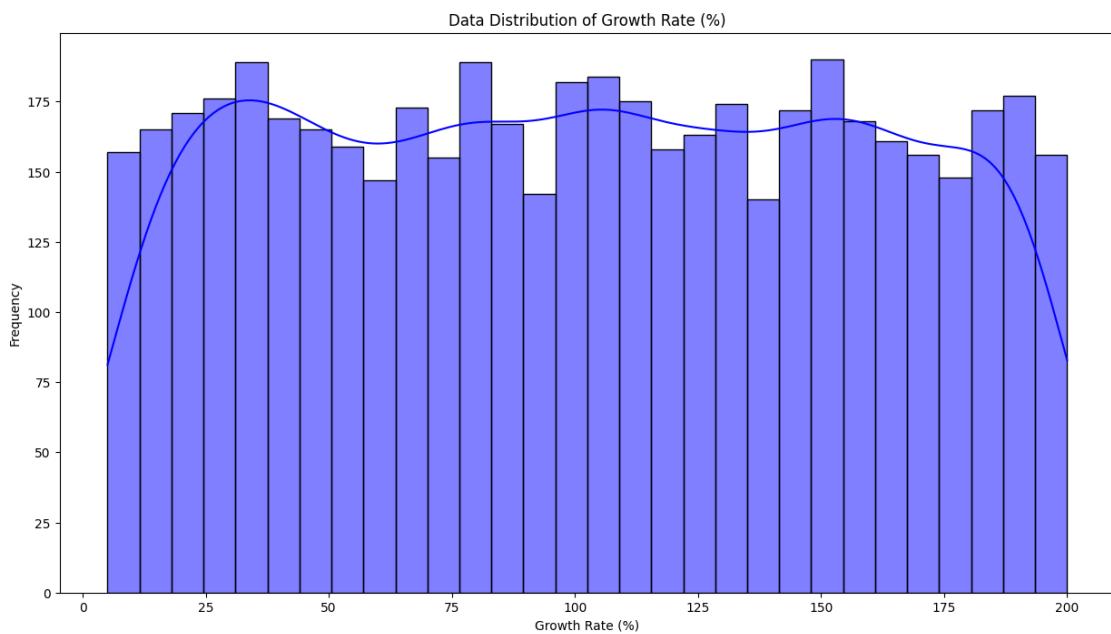
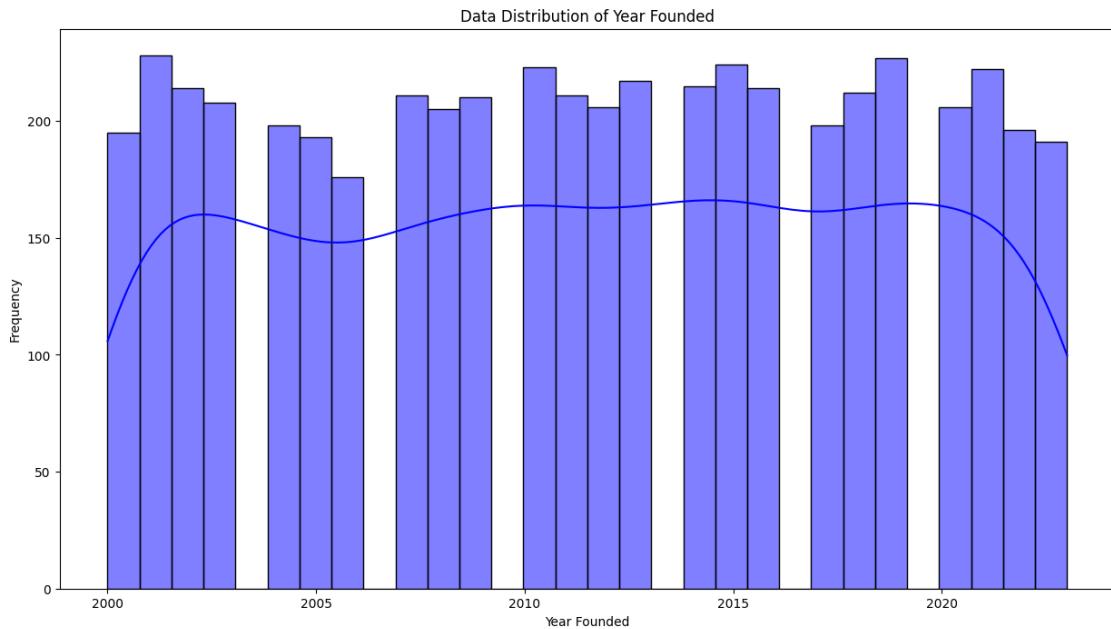
```

[11]: for col in df.select_dtypes(include=['int64', 'float64']).columns:
    plt.figure(figsize=(15, 8))
    sns.histplot(df[col], kde=True, color='blue', bins=30)
    plt.title(f"Data Distribution of {col}")
    plt.ylabel("Frequency")
    plt.xlabel(col)
    plt.show()

```







```
[12]: # Select numerical columns (excluding the target variable)
df_categorical = df.select_dtypes(include=['object', 'int64']).  
    ↪drop(columns=['Growth Rate (%)'], errors='ignore')
```

```
# Get the number of numerical columns
```

```

col_count = len(df_categorical.columns)
nrows = (col_count + 1) // 2 # Ensure enough rows for subplots
ncols = 2

fig, axes = plt.subplots(nrows=nrows, ncols=ncols, figsize=(15, 5 * nrows))
axes = axes.flatten() if col_count > 1 else [axes]

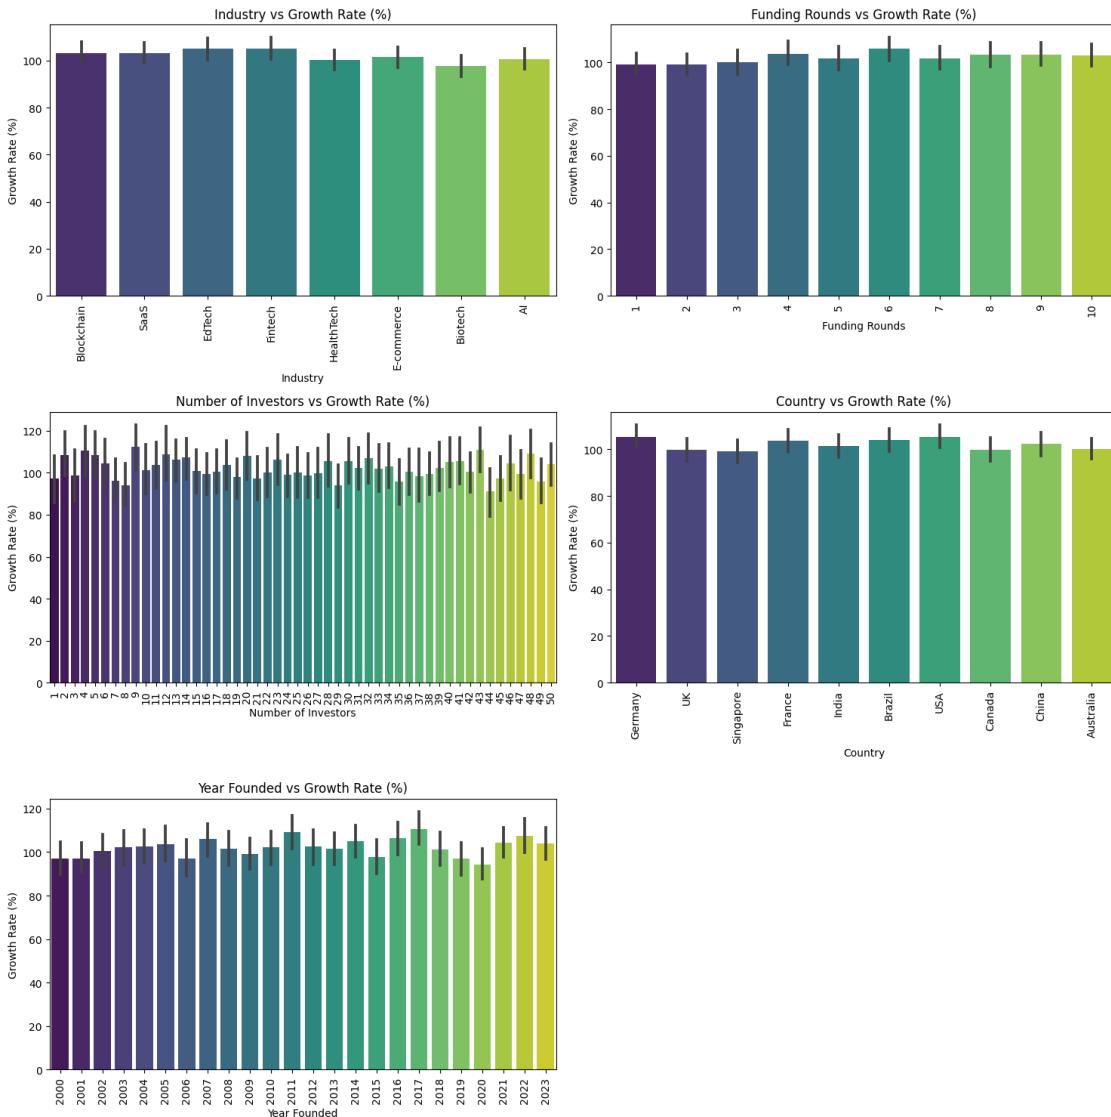
# Plot each numerical variable against Early_Prediction
for i, column in enumerate(df_categorical.columns):
    sns.barplot(x=df[column], y=df['Growth Rate (%)'], ax=axes[i],  

    ↪ palette='viridis')
    axes[i].set_title(f'{column} vs Growth Rate (%)')
    axes[i].tick_params(axis='x', rotation=90)

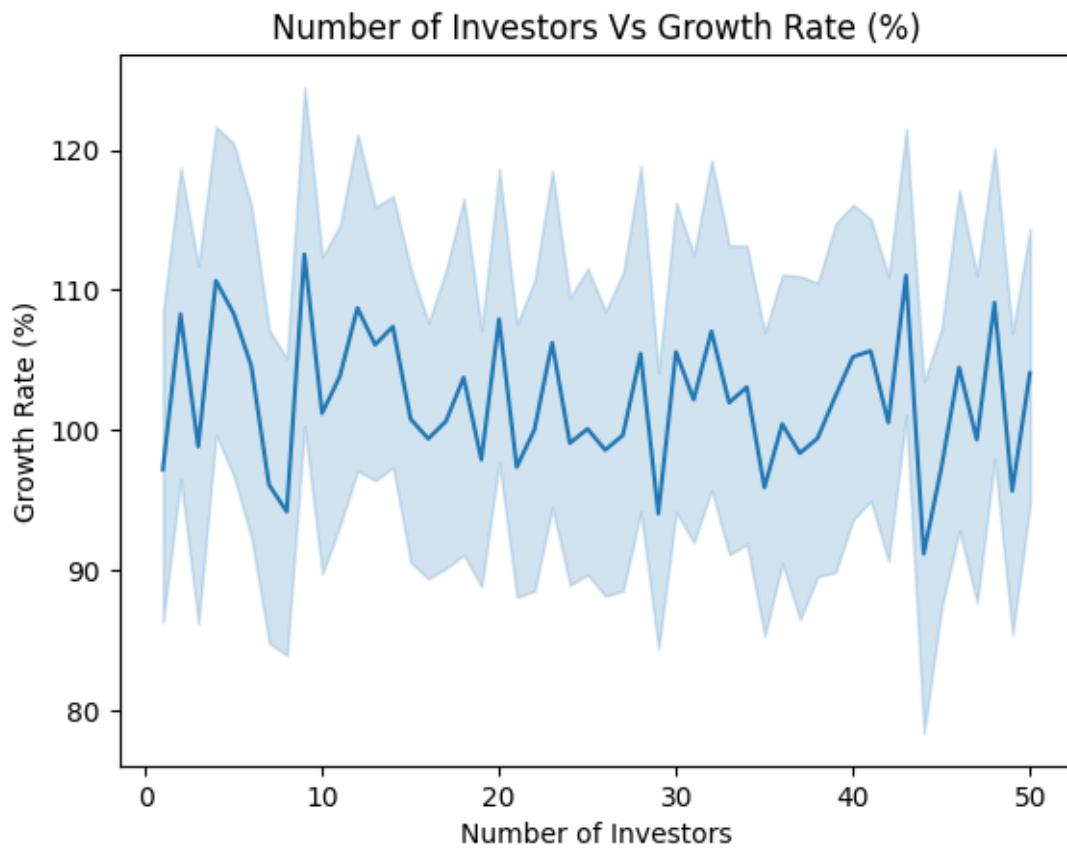
# Remove unused subplots if the number of numerical columns is odd
if col_count % 2 != 0:
    fig.delaxes(axes[-1])

plt.tight_layout()
plt.show()

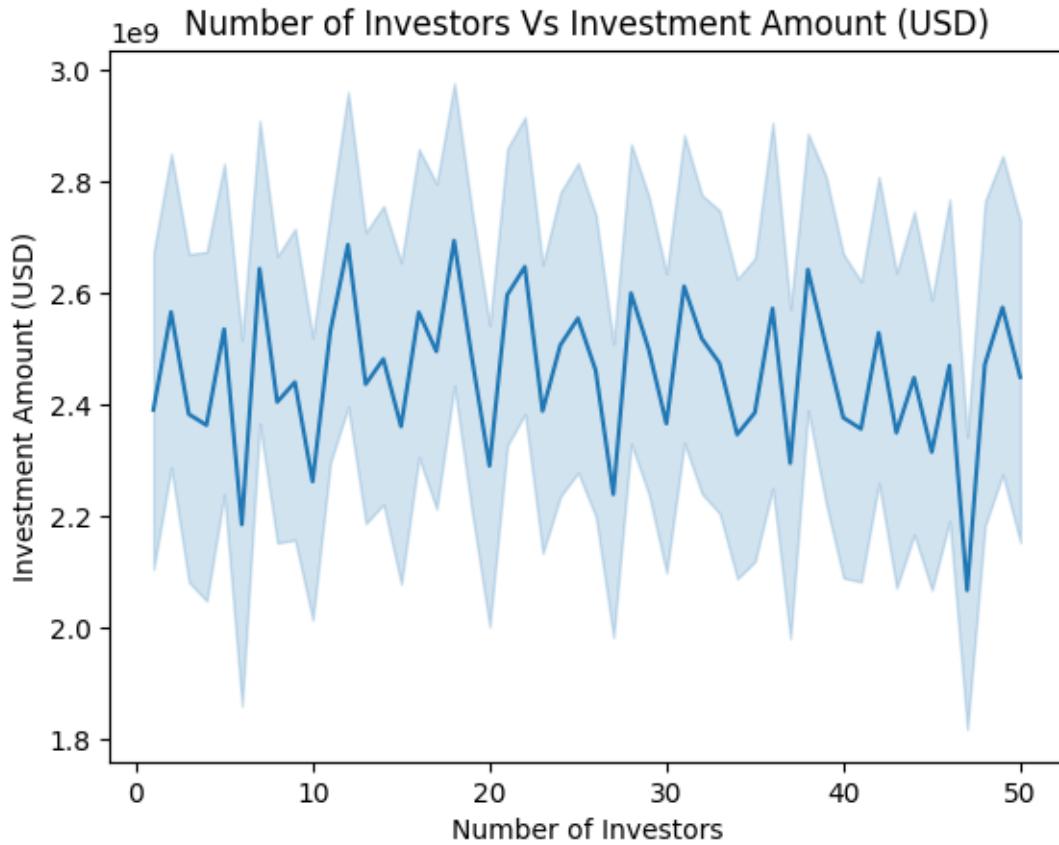
```



```
[13]: sns.lineplot(x="Number of Investors", y="Growth Rate (%)", data=df)
plt.title("Number of Investors Vs Growth Rate (%)")
plt.show()
```



```
[14]: sns.lineplot(x="Number of Investors", y="Investment Amount (USD)", data=df)
plt.title("Number of Investors Vs Investment Amount (USD)")
plt.show()
```



3 Machine Learning

Preprocessing Data

```
[15]: from sklearn.preprocessing import LabelEncoder

# Select categorical columns
df_categorical = df.select_dtypes(include='object')

# Initialize LabelEncoder
label_encoders = {}

# Apply LabelEncoder to each categorical column
for column in df_categorical.columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column]) # Transform the categorical ↵column
    label_encoders[column] = le # Store the encoder for future inverse ↵transformation if needed
```

```
# Check the transformed dataframe
print(df.head())
```

	Industry	Funding Rounds	Investment Amount (USD)	Valuation (USD)	\
0	2	8	1.335166e+09	6.621448e+09	
1	7	2	2.781498e+09	8.363214e+09	
2	4	10	3.309032e+09	1.548270e+10	
3	5	5	4.050196e+09	1.268253e+10	
4	4	9	1.645080e+09	6.887966e+09	

	Number of Investors	Country	Year Founded	Growth Rate (%)
0	50	5	2012	77.10
1	36	8	2006	105.52
2	39	7	2016	190.47
3	44	4	2021	9.44
4	48	6	2011	192.00

Loading the training and testing data

```
[16]: from sklearn.model_selection import train_test_split
import numpy as np

X = df.drop("Growth Rate (%)", axis=1)
Y = np.log1p(df["Growth Rate (%")]

# Scale Fit Transformation
X_scaled = StandardScaler().fit_transform(X)
```

```
[17]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, Y, test_size=0.2,
random_state=42)
```

Setting up and training the best model

```
[18]: import joblib
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor,
AdaBoostRegressor
from sklearn.svm import SVR

# Split dataset into training and testing (assuming X_scaled and Y are already
defined)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, Y, test_size=0.2,
random_state=42)
```

```

# Define parameter grids for each model
param_grids = {
    "Linear Regression": {
        "fit_intercept": [True, False]
    },
    "Ridge Regression": {
        "alpha": [0.1, 1, 10, 50, 100]
    },
    "Lasso Regression": {
        "alpha": [0.001, 0.01, 0.1, 1, 10]
    },
    "ElasticNet Regression": {
        "alpha": [0.001, 0.01, 0.1, 1, 10],
        "l1_ratio": [0.1, 0.5, 0.9]
    },
    "Decision Tree": {
        "max_depth": [None, 5, 10, 20],
        "min_samples_split": [2, 5, 10]
    },
    "Random Forest": {
        "n_estimators": [100, 200],
        "max_depth": [None, 5, 10],
        "min_samples_split": [2, 5]
    },
    "Gradient Boosting": {
        "n_estimators": [100, 200],
        "learning_rate": [0.01, 0.1, 0.2],
        "max_depth": [3, 5, 7]
    },
    "AdaBoost": {
        "n_estimators": [50, 100, 200],
        "learning_rate": [0.01, 0.1, 1]
    },
    "Support Vector Machine": {
        "C": [0.1, 1, 10],
        "epsilon": [0.01, 0.1, 0.2],
        "kernel": ['rbf']
    }
}

# Define base models
models = {
    "Linear Regression": LinearRegression(),
    "Ridge Regression": Ridge(random_state=42),
    "Lasso Regression": Lasso(random_state=42),
    "ElasticNet Regression": ElasticNet(random_state=42),
}

```

```

    "Decision Tree": DecisionTreeRegressor(random_state=42),
    "Random Forest": RandomForestRegressor(random_state=42),
    "Gradient Boosting": GradientBoostingRegressor(random_state=42),
    "AdaBoost": AdaBoostRegressor(random_state=42),
    "Support Vector Machine": SVR()
}

best_model = None
best_r2 = float('-inf')
results = []

# Tune hyperparameters for each model using GridSearchCV
for name, model in models.items():
    print(f"Tuning {name}...")
    grid = GridSearchCV(model, param_grid=param_grids[name], cv=5, u
    ↪scoring='r2', n_jobs=-1)
    grid.fit(X_train, y_train)
    best_estimator = grid.best_estimator_
    predictions = best_estimator.predict(X_test)

    # Calculate evaluation metrics
    r2 = r2_score(y_test, predictions)
    mse = mean_squared_error(y_test, predictions)
    mae = mean_absolute_error(y_test, predictions)

    results.append((name, r2, mse, mae, grid.best_params_))

    print(f"Best params for {name}: {grid.best_params_}")
    print(f"R² Score: {r2:.4f}, MSE: {mse:.4f}, MAE: {mae:.4f}\n")

    # Track the best model by R² score
    if r2 > best_r2:
        best_r2 = r2
        best_model = (name, best_estimator)

# Print a summary of model performances
print("\nModel Performance Comparison:")
print("=" * 40)
for name, r2, mse, mae, best_params in results:
    print(f"{name}:")
    print(f"    R² Score: {r2:.4f}")
    print(f"    Mean Squared Error: {mse:.4f}")
    print(f"    Mean Absolute Error: {mae:.4f}")
    print(f"    Best Params: {best_params}")
    print("-" * 40)

# Save the best model to disk

```

```
best_model_name, best_model_instance = best_model
joblib.dump(best_model_instance, "growth_rate_startup_prediction_model.pkl")
print(f"\n Best Model: {best_model_name} with R2 Score: {best_r2:.4f} (Saved as 'best_model.pkl')")
```

Tuning Linear Regression...

Best params for Linear Regression: {'fit_intercept': True}
R² Score: -0.0024, MSE: 0.5955, MAE: 0.6275

Tuning Ridge Regression...

Best params for Ridge Regression: {'alpha': 100}
R² Score: -0.0021, MSE: 0.5954, MAE: 0.6274

Tuning Lasso Regression...

Best params for Lasso Regression: {'alpha': 0.1}
R² Score: -0.0007, MSE: 0.5945, MAE: 0.6269

Tuning ElasticNet Regression...

Best params for ElasticNet Regression: {'alpha': 0.1, 'l1_ratio': 0.5}
R² Score: -0.0007, MSE: 0.5945, MAE: 0.6269

Tuning Decision Tree...

Best params for Decision Tree: {'max_depth': 5, 'min_samples_split': 2}
R² Score: -0.0093, MSE: 0.5996, MAE: 0.6283

Tuning Random Forest...

Best params for Random Forest: {'max_depth': 5, 'min_samples_split': 5, 'n_estimators': 100}
R² Score: -0.0059, MSE: 0.5976, MAE: 0.6282

Tuning Gradient Boosting...

Best params for Gradient Boosting: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}
R² Score: 0.0004, MSE: 0.5939, MAE: 0.6259

Tuning AdaBoost...

Best params for AdaBoost: {'learning_rate': 0.01, 'n_estimators': 50}
R² Score: -0.0045, MSE: 0.5968, MAE: 0.6317

Tuning Support Vector Machine...

Best params for Support Vector Machine: {'C': 0.1, 'epsilon': 0.2, 'kernel': 'rbf'}
R² Score: -0.0788, MSE: 0.6410, MAE: 0.6043

Model Performance Comparison:

=====

Linear Regression:

```
R2 Score: -0.0024
Mean Squared Error: 0.5955
Mean Absolute Error: 0.6275
Best Params: {'fit_intercept': True}

-----
Ridge Regression:
R2 Score: -0.0021
Mean Squared Error: 0.5954
Mean Absolute Error: 0.6274
Best Params: {'alpha': 100}

-----
Lasso Regression:
R2 Score: -0.0007
Mean Squared Error: 0.5945
Mean Absolute Error: 0.6269
Best Params: {'alpha': 0.1}

-----
ElasticNet Regression:
R2 Score: -0.0007
Mean Squared Error: 0.5945
Mean Absolute Error: 0.6269
Best Params: {'alpha': 0.1, 'l1_ratio': 0.5}

-----
Decision Tree:
R2 Score: -0.0093
Mean Squared Error: 0.5996
Mean Absolute Error: 0.6283
Best Params: {'max_depth': 5, 'min_samples_split': 2}

-----
Random Forest:
R2 Score: -0.0059
Mean Squared Error: 0.5976
Mean Absolute Error: 0.6282
Best Params: {'max_depth': 5, 'min_samples_split': 5, 'n_estimators': 100}

-----
Gradient Boosting:
R2 Score: 0.0004
Mean Squared Error: 0.5939
Mean Absolute Error: 0.6259
Best Params: {'learning_rate': 0.01, 'max_depth': 3, 'n_estimators': 100}

-----
AdaBoost:
R2 Score: -0.0045
Mean Squared Error: 0.5968
Mean Absolute Error: 0.6317
Best Params: {'learning_rate': 0.01, 'n_estimators': 50}

-----
Support Vector Machine:
```

```
R2 Score: -0.0788
Mean Squared Error: 0.6410
Mean Absolute Error: 0.6043
Best Params: {'C': 0.1, 'epsilon': 0.2, 'kernel': 'rbf'}
```

```
Best Model: Gradient Boosting with R2 Score: 0.0004 (Saved as
'best_model.pkl')
```

4 Conclusion

- There is a very strong positive relationship between the valuation of a startup and the total investment put into that startup, indicating that the amount a startup is valued at is correlates highly with the total investment put into it.
- The Industry with the highest growth rate is **EdTech**
- The Funding rounds with the highest average **Growth Rate** is **6**
- The Countries with the highest Growth Rate is **Germany and USA**
- The Average Growth Rate according to the dataset is **102.1%**
- There is little to no relationship between the **Growth Rate** of a startup and all the other factors in the dataset like the number of investors, the year the business was founded, the total investment amount or the business' valuation.
- The Gradient Boosting model performed moderately okay, with an R² Score of **0.0004**.

Ways to increase Growth Rate may include - Product Diversification – Expand your offerings to attract more customers and reduce reliance on a single revenue source. - Market Expansion – Enter new regions or target different demographics to increase your customer base. - Strategic Acquisitions – Acquire other businesses to gain new markets, technologies, and competitive advantages.